

Learning When to Delay, Abstain, or Act: Risk-Aware Decision Policies Under Uncertainty

Rajveer Singh Arneja
Email: rajveer.arneja779@nmims.in

Vaibhavee Tiwari
Email: vaibhavee.tiwari780@nmims.in

Abstract—Machine learning systems are increasingly deployed in environments where incorrect decisions carry asymmetric and often irreversible consequences. While selective prediction methods allow models to abstain under low confidence, they assume that decisions must either be made immediately or rejected outright. In many operational settings, however, actions can be deferred to gather additional evidence before committing to a prediction.

This paper studies a delay-aware decision framework in which a model can choose among three actions: predict, abstain, or defer. We formalize this problem as a risk-sensitive policy learning task that explicitly accounts for uncertainty, cost asymmetry, and potential information gain over time. Unlike standard classification pipelines, the proposed approach evaluates decisions not only by predictive accuracy but by expected operational cost under distributional variability.

Through controlled experiments on fraud detection and simulated monitoring environments, we demonstrate that incorporating delay as an explicit policy option reduces high-cost errors under distribution shift. The results highlight a practical trade-off between responsiveness and reliability, suggesting that sequential evidence acquisition can significantly improve robustness in safety-critical systems.

I. INTRODUCTION

Machine learning models are increasingly deployed in decision pipelines where incorrect predictions incur asymmetric and potentially irreversible costs. Examples include fraud detection systems that block legitimate transactions, monitoring systems that prematurely escalate incidents, and automated approval systems that reject valid applications. In such settings, prediction accuracy alone is insufficient as a measure of system reliability. The consequences of errors vary in magnitude, timing, and operational impact.

Conventional classification models are designed to produce immediate predictions once an input is observed. This design implicitly assumes that the decision must be made at a single point in time. However, real-world decision processes are often sequential. In many operational environments, systems can defer action, gather additional information, or continue monitoring before committing to a final decision. Despite this practical reality, most machine learning pipelines treat prediction as an instantaneous event rather than a temporally adaptive process.

Selective prediction has emerged as one mechanism to improve reliability. By allowing models to abstain when predictive confidence falls below a threshold, selective classifiers reduce the risk of high-confidence errors. While abstention improves safety in uncertain cases, it assumes that declining

to predict is the only alternative to immediate action. In many systems, however, abstention and delay are fundamentally different. Abstention may transfer responsibility to a human operator, whereas delay may allow automated systems to accumulate evidence before deciding.

The distinction between abstaining and delaying becomes particularly important under distribution shift. When deployment data differs from training data, model confidence estimates may become unreliable. In such cases, committing to early predictions can amplify risk. A delay mechanism introduces the possibility of revising uncertainty estimates as new signals arrive, potentially reducing high-cost mistakes without permanently rejecting the input.

This paper investigates a unified decision framework in which a machine learning system may choose among three actions: act immediately, abstain from prediction, or delay the decision to gather additional evidence. We formalize this as a risk-aware policy learning problem that explicitly models asymmetric costs and uncertainty dynamics. Rather than optimizing for accuracy alone, the objective is to minimize expected operational risk under both nominal and shifted distributions.

Our approach differs from traditional selective prediction in two key ways. First, we introduce delay as an explicit action within the decision space, allowing sequential refinement of confidence estimates. Second, we evaluate policies under controlled distribution shift to study robustness beyond static validation settings. The goal is not to design a new predictive model architecture, but to understand how decision policies can be structured to improve reliability in realistic environments.

We evaluate the proposed framework using both real-world fraud detection data and simulated monitoring environments that reflect incident escalation dynamics. Across experiments, we analyze trade-offs between responsiveness, coverage, and expected cost. The results suggest that incorporating delay as an intermediate decision state significantly reduces high-cost errors in uncertain regimes, particularly under moderate distributional variability.

The contributions of this work are summarized as follows:

- We formalize delay-aware decision-making as a tri-action policy problem integrating prediction, abstention, and deferral.
- We introduce a risk-sensitive evaluation framework that captures asymmetric costs and temporal information gain.

- We provide empirical analysis demonstrating improved robustness under distribution shift through delay-aware policies.
- We analyze trade-offs between decision latency and reliability in operational contexts.

By reframing classification as a sequential decision problem rather than a one-shot prediction task, this work aims to bridge the gap between machine learning theory and deployment-level reliability requirements.

II. RELATED WORK

Improving reliability of machine learning decisions has been studied across several research directions, including selective prediction, uncertainty estimation, cost-sensitive learning, and distribution shift robustness. These areas provide important foundations for designing systems that operate safely under imperfect information.

Selective prediction introduces the ability for models to abstain from making predictions when confidence is insufficient. The classical formulation of classification with a reject option dates back to early decision-theoretic work [1]. More recent approaches extend selective classification to modern deep learning systems by optimizing risk–coverage trade-offs and leveraging calibrated confidence estimates [2], [3]. These methods typically define a confidence threshold below which predictions are withheld. While effective in reducing unreliable outputs, selective prediction frameworks generally assume that abstention is the only alternative to immediate decision-making.

Uncertainty estimation methods aim to quantify predictive confidence through probabilistic modeling, Bayesian approximations, or ensemble techniques. Deep ensembles [4] and Bayesian neural approximations [5] have demonstrated improved uncertainty estimates compared to deterministic models. Confidence calibration has also been widely studied as a mechanism to improve reliability of predicted probabilities, particularly under distribution shift [6]. Despite these advances, most uncertainty-aware approaches still operate within a static decision paradigm in which each input must be classified or rejected at a single time step.

Risk-sensitive and cost-sensitive learning frameworks address asymmetric consequences of prediction errors by incorporating application-specific loss functions into training objectives. In fraud detection, medical diagnosis, and monitoring systems, false positives and false negatives often carry significantly different operational costs. Cost-sensitive classifiers therefore adjust decision boundaries to minimize expected loss rather than classification error. However, these models typically assume that decisions must be made immediately once an input is observed, without considering the possibility of postponing action to gather additional evidence.

Sequential decision-making and reinforcement learning literature provides theoretical tools for modeling decisions that evolve over time. These frameworks analyze policies that balance action cost, information acquisition, and long-term reward optimization. While sequential decision models capture

temporal dynamics, they are rarely integrated with selective prediction mechanisms in supervised classification settings. In many practical pipelines, prediction remains a one-shot process despite evidence that decision reliability can benefit from sequential evidence accumulation.

Research on distribution shift further highlights the limitations of static decision systems. Dataset shift and domain generalization studies show that models trained under stationary assumptions often degrade when deployed in evolving environments [7]. Out-of-distribution detection methods attempt to identify uncertain inputs [8], but typically focus on flagging anomalies rather than adapting decision timing as a risk mitigation strategy.

The present work builds upon these research directions by introducing delay as an explicit action alongside prediction and abstention. By integrating selective prediction with sequential decision reasoning, the proposed framework extends reliability-focused learning beyond static classification settings and enables adaptive decision timing under uncertainty and cost asymmetry.

III. PROBLEM FORMULATION

Consider a supervised classification setting in which an input instance $x \in \mathcal{X}$ is associated with a ground-truth label $y \in \mathcal{Y}$. A predictive model f produces a posterior probability distribution over possible labels:

$$p(y|x) = f(x) \quad (1)$$

Traditional classification systems generate an immediate prediction by selecting the most probable label:

$$\hat{y} = \arg \max_{y \in \mathcal{Y}} p(y|x) \quad (2)$$

This formulation implicitly assumes that decisions must be made at the moment an input is observed. However, in many real-world decision environments, actions may be sequential and can benefit from additional information gathered over time. Examples include fraud monitoring systems, safety-critical alerts, and automated approval pipelines, where premature decisions may incur significant operational risk.

Selective prediction frameworks extend standard classification by allowing models to abstain from prediction when confidence is insufficient. Formally, abstention introduces a rejection action when predictive confidence falls below a predefined threshold τ . While abstention reduces unreliable predictions, it assumes that declining to decide is the only alternative to immediate action.

In contrast, many practical systems allow decisions to be deferred rather than permanently rejected. To capture this behavior, we extend the selective prediction framework by introducing a delay action that enables the system to postpone decisions while acquiring additional evidence. The resulting decision policy operates over an action space defined as:

$$\mathcal{A} = \{a_{pred}, a_{abs}, a_{delay}\} \quad (3)$$

where a_{pred} denotes immediate prediction, a_{abs} represents abstention, and a_{delay} corresponds to deferring the decision.

A. Risk-Aware Decision Objective

Decision-making in safety-critical and cost-sensitive environments involves asymmetric penalties associated with different actions. To model these trade-offs, we define a risk-aware objective that evaluates the cost associated with each possible decision outcome.

Let $a \in \mathcal{A}$ denote the selected action for input x . The action-dependent risk function is defined as:

$$R(a, x) = \begin{cases} C_{err} \cdot \mathbb{I}[\hat{y} \neq y], & \text{if } a = a_{pred} \\ C_{abs}, & \text{if } a = a_{abs} \\ C_{delay}(\Delta t) - \gamma I(x, \Delta t), & \text{if } a = a_{delay} \end{cases} \quad (4)$$

Here, C_{err} denotes the cost associated with incorrect predictions, C_{abs} represents the operational overhead introduced by abstention, and $C_{delay}(\Delta t)$ models the cost incurred by postponing decisions over time interval Δt . The term $I(x, \Delta t)$ represents the expected information gain obtained during the delay period, while γ controls the relative importance assigned to uncertainty reduction.

The optimal decision policy selects the action that minimizes expected risk:

$$a^*(x) = \arg \min_{a \in \mathcal{A}} \mathbb{E}[R(a, x)] \quad (5)$$

This formulation enables the decision system to jointly balance predictive accuracy, decision latency, and operational cost within a unified optimization framework. By incorporating delay as an explicit decision option, the proposed formulation extends conventional selective prediction into a sequential decision-making paradigm that is better aligned with real-world deployment constraints.

IV. DELAY-AWARE DECISION FRAMEWORK

The proposed framework evaluates candidate actions for each input instance using confidence-aware decision rules. For each observation x , the system may choose to produce an immediate prediction, abstain from decision, or delay action to collect additional information.

A. Confidence Estimation

Prediction confidence is estimated using calibrated posterior probabilities:

$$Conf(x) = \max_{y \in \mathcal{Y}} p(y|x) \quad (6)$$

Calibration is applied to improve reliability of probability estimates under distributional variability.

B. Tri-Action Decision Policy

The decision policy is parameterized using two confidence thresholds, τ_{pred} and τ_{delay} , which partition the confidence space into three operational regimes. The action selection rule is defined as:

$$a(x) = \begin{cases} a_{pred}, & Conf(x) \geq \tau_{pred} \\ a_{delay}, & \tau_{delay} \leq Conf(x) < \tau_{pred} \\ a_{abs}, & Conf(x) < \tau_{delay} \end{cases} \quad (7)$$

This tri-action policy generalizes conventional selective prediction by incorporating temporal flexibility into decision-making.

V. DELAY-AWARE DECISION ALGORITHM

The proposed tri-action policy is implemented using confidence-based threshold evaluation combined with risk-aware decision rules. The algorithm evaluates candidate actions for each input instance and selects the action minimizing expected operational risk.

Algorithm 1 Tri-Action Delay-Aware Decision Policy

Require: Input instance x , trained model f , thresholds τ_{pred} ,

τ_{delay}

- 1: Compute probability distribution $p(y|x) = f(x)$
 - 2: Compute confidence score $Conf(x) = \max_y p(y|x)$
 - 3: **if** $Conf(x) \geq \tau_{pred}$ **then**
 - 4: Output predicted label
 - 5: **else if** $\tau_{delay} \leq Conf(x) < \tau_{pred}$ **then**
 - 6: Delay decision and acquire additional evidence
 - 7: **else**
 - 8: Abstain from prediction
 - 9: **end if**
-

Threshold parameters were selected using validation risk minimization.

VI. EXPERIMENTAL SETUP

A. Datasets

We evaluate the proposed framework using the publicly available Credit Card Fraud Detection dataset, which contains anonymized European card transactions collected over a two-day period. The dataset exhibits severe class imbalance, with fraudulent transactions accounting for less than 0.2% of total observations, making it suitable for studying uncertainty-aware and cost-sensitive decision policies.

To further evaluate robustness under evolving environments, a synthetic monitoring setting was created through controlled perturbations including class imbalance variation, feature noise injection, and confidence degradation. These settings simulate deployment conditions where real-world distributions deviate from training assumptions.

B. Models

Experiments employ probabilistic classifiers capable of producing calibrated confidence estimates, including logistic regression and random forest models. These models are selected to isolate decision policy behavior without introducing architectural complexity.

C. Distribution Shift Simulation

To evaluate robustness, controlled distribution shifts are introduced through class imbalance variation, feature perturbation, and noise injection. These shifts simulate realistic deployment scenarios where data distributions evolve over time.

VII. EVALUATION METRICS

Performance is evaluated using metrics that capture reliability and operational cost trade-offs. These include selective risk, prediction coverage, expected operational cost, and calibration error. Together, these measures provide a comprehensive assessment of decision quality beyond conventional accuracy metrics.

VIII. RESULTS

Experimental evaluation demonstrates that incorporating delay as an explicit decision action improves reliability under uncertainty. Selective prediction experiments reveal that abstention policies reduce prediction error in low-confidence regions, but delay-aware policies provide additional flexibility by allowing the system to refine uncertainty estimates before committing to decisions.

Risk-coverage analysis shows that delay-aware policies achieve lower expected operational cost compared to standard selective prediction across a wide range of confidence thresholds. Under simulated distribution shift, models relying solely on immediate prediction exhibit noticeable degradation in reliability, particularly in regions where confidence calibration becomes unstable. Introducing delay reduces the frequency of high-cost misclassifications by enabling additional evidence accumulation prior to decision commitment.

Cost-sensitive evaluation further indicates that delay policies outperform abstention-only strategies when the relative cost of incorrect automated decisions significantly exceeds the operational cost of deferral. Across both fraud detection and monitoring simulation experiments, tri-action decision policies demonstrate improved robustness while maintaining acceptable prediction coverage.

A. Model Comparison

To evaluate the generality of the proposed decision framework, experiments were conducted using multiple probabilistic classifiers, including logistic regression and random forest models. The objective of this comparison is to analyze whether delay-aware policies provide consistent reliability improvements across different model families.

Results indicate that both models benefit from selective decision policies, although performance characteristics differ.




Fig. 1: Risk-coverage trade-off for the Random Forest model. Increasing selectivity reduces prediction risk while sacrificing automated coverage.




Fig. 2: Expected operational cost comparison across immediate prediction, abstention-only, and tri-action decision policies. Lower cost indicates improved reliability under uncertainty.

Logistic regression exhibits smoother confidence threshold transitions, while random forest models achieve stronger reliability performance under uncertainty.

For logistic regression models, the tri-action strategy achieved the lowest selective risk (0.0812) and highest selective F1 performance relative to baseline prediction. For random forest classifiers, cost-aware selective prediction achieved the lowest selective risk (0.0133) and expected operational cost (0.1169), indicating superior reliability under uncertainty. However, the tri-action framework preserved substantially higher automated coverage (96.6%) by routing uncertain predictions through a deferral mechanism rather than abstaining completely.

These findings suggest a practical trade-off between maximizing reliability and preserving automation coverage in operational systems.

B. Distribution Shift Analysis

To evaluate robustness, controlled distribution shifts were introduced through class imbalance variation and feature

TABLE I: Performance comparison across selective prediction strategies. Lower selective risk and cost are preferred.

Model	Strategy	Coverage	Risk	Cost
LR	Baseline	1.000	0.1915	0.2235
LR	Static	0.699	0.1030	0.2385
LR	Dynamic	0.800	0.1266	0.2222
LR	Tri-Action	0.6158	0.0812	0.5639
RF	Baseline	1.000	0.0467	0.2328
RF	Static	0.9782	0.0350	0.1821
RF	Cost-Aware	0.8818	0.0133	0.1169
RF	Tri-Action	0.9660	0.0305	0.1955




Fig. 3: Selective risk under distribution shift for the Random Forest model. Delay-aware decision strategies maintain stronger reliability under increasing distributional variability.

noise injection. These shifts simulate deployment environments where real-world data distributions deviate from training conditions.

Under distribution shift, baseline models relying on immediate prediction exhibit increased misclassification rates, particularly in low-confidence regions. Selective prediction mitigates this degradation by rejecting uncertain predictions; however, delay-aware policies provide additional protection by allowing the model to refine confidence estimates over time.

Experimental results demonstrate that delay-based policies maintain lower expected operational cost under moderate distributional variability. This behavior highlights the effectiveness of temporal decision flexibility as a reliability enhancement mechanism when confidence calibration alone becomes insufficient.

IX. DISCUSSION

The results highlight the importance of incorporating temporal flexibility into machine learning decision pipelines. Conventional classification systems treat predictions as instantaneous events, whereas many real-world decision environments naturally support sequential evidence acquisition. The proposed delay-aware framework provides a structured mechanism to leverage this property.

One key observation is the trade-off between responsiveness and reliability. While delay reduces catastrophic prediction

errors, excessive deferral may increase operational latency and reduce system throughput. The optimal delay threshold therefore depends on application-specific cost structures and decision urgency. In high-risk environments such as fraud detection or safety monitoring, conservative policies that prioritize reliability may be preferable.

Another important finding is the interaction between distribution shift and confidence calibration. Under moderate distributional variability, confidence scores alone may not reliably indicate prediction correctness. Delay-based policies provide an additional safeguard by allowing uncertainty to be reassessed over time. This suggests that temporal decision strategies may serve as complementary mechanisms to conventional calibration methods.

The proposed framework remains intentionally model-agnostic, allowing integration with a wide range of probabilistic classifiers. However, the current implementation assumes simplified information gain estimation during delay intervals. More sophisticated modeling of temporal signal dynamics could further improve decision quality.

A. Limitations and Practical Deployment Considerations

While the proposed framework improves reliability, several limitations remain. The current delay modeling assumes simplified estimation of information gain during deferral intervals. In real-world systems, information acquisition dynamics may depend on complex temporal signal behavior that requires sequential modeling techniques.

Operational deployment also introduces constraints related to decision latency and system throughput. Excessive deferral may negatively impact user experience or operational efficiency. Therefore, threshold selection must balance reliability improvement against latency tolerance specific to the application domain.

Additionally, the framework assumes availability of calibrated confidence estimates. In environments with severe distribution shift or adversarial conditions, calibration reliability may degrade, potentially affecting decision stability.

Future work should investigate adaptive threshold learning and reinforcement learning-based delay policy optimization to address these challenges.

X. REPRODUCIBILITY AND IMPLEMENTATION

To support reproducibility and transparency, implementation code, dataset preprocessing scripts, and experimental configurations used in this study are publicly available in an open-source repository:

<https://github.com/rajveersingharneja/selective-prediction-uncertainty>

The repository includes model training pipelines, selective decision policy implementations, and scripts used to generate evaluation plots presented in this paper.

XI. CONCLUSION

This work presents a unified framework for tri-action decision-making that extends conventional selective prediction

by incorporating delay as an explicit operational strategy. By jointly modeling prediction, abstention, and deferral within a risk-aware formulation, the proposed approach enables machine learning systems to adapt decision timing based on uncertainty estimates and asymmetric cost considerations.

Experimental evaluation demonstrates that delay-aware policies reduce high-cost prediction errors under distribution shift while preserving practical coverage levels. The results suggest that incorporating temporal flexibility into classification pipelines improves reliability in environments where immediate decisions may be unsafe or suboptimal. In particular, the proposed framework highlights the value of sequential evidence accumulation as a complementary mechanism to traditional confidence-based rejection strategies.

Future work will explore adaptive delay policy learning using reinforcement learning, improved estimation of temporal information gain, and deployment of delay-aware decision frameworks in streaming and real-time monitoring systems. These directions offer promising opportunities for improving the safety and robustness of machine learning systems operating in dynamic environments.

REFERENCES

- [1] C. Chow, “On optimum error and reject tradeoff,” *IEEE Transactions on Information Theory*, vol. 16, no. 1, pp. 41–46, 1970.
- [2] Y. Geifman and R. El-Yaniv, “Selective classification for deep neural networks,” *Advances in Neural Information Processing Systems*, 2017.
- [3] R. El-Yaniv, “Foundations of selective classification,” *Journal of Machine Learning Research*, vol. 11, pp. 1605–1641, 2010.
- [4] B. Lakshminarayanan, A. Pritzel, and C. Blundell, “Simple and scalable predictive uncertainty estimation using deep ensembles,” in *Advances in Neural Information Processing Systems*, 2017.
- [5] Y. Gal and Z. Ghahramani, “Dropout as a bayesian approximation: Representing model uncertainty in deep learning,” in *International Conference on Machine Learning*, 2016.
- [6] C. Guo, G. Pleiss, Y. Sun, and K. Weinberger, “On calibration of modern neural networks,” in *International Conference on Machine Learning*, 2017.
- [7] J. Quinonero-Candela, M. Sugiyama, A. Schwaighofer, and N. Lawrence, *Dataset Shift in Machine Learning*. MIT Press, 2009.
- [8] D. Hendrycks and K. Gimpel, “A baseline for detecting misclassified and out-of-distribution examples,” in *International Conference on Learning Representations*, 2017.